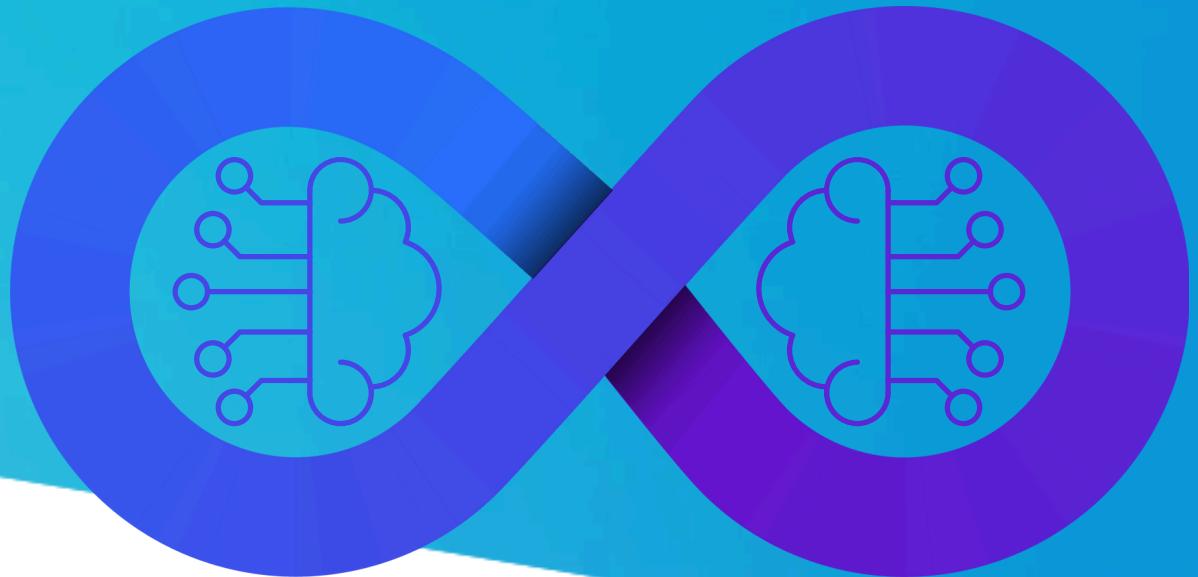




MLOps with Agentic AI

Advanced Certification Program

Curriculum





Index

Module 1: Python and MLOps Foundations

- Python + Numpy + Pandas
- MLOps overview, CI/CD for ML, MLflow, model versioning

Module 2: Modern Cloud-Native MLOps

- Airflow, Kubeflow, Optuna
- Kubernetes, TF Serving, Prometheus, Grafana
- Advanced deployment (blue-green, canary)
- Feedback loops, retraining, and cost optimization

Module 3: Cloud & Productionization

- AWS pipelines with SageMaker/ML Studio
- Model registry, approval workflows, drift monitoring
- Multi-cloud pipeline integration
- Capstone 1: End-to-end MLOps pipeline

Module 4: Agentic AI & LLMOps

- LLM lifecycle, Fine-tuning vs Prompt-tuning
- RAG architecture, vector DBs
- Deploying LLMs (Kubernetes, LangServe)
- LangChain, CrewAI, LangGraph
- Evaluation tools: RAGAS, LangSmith, Trulens
- Capstone 2: Build & deploy an agentic AI system for enterprise use-case

1 Python and MLOps Foundations

Session 1: Python Refresher for MLOps (Beginner)

Goal: Develop modular Python code structures essential for machine learning engineering, emphasizing reusability, logging, configuration, and maintainability.

Topics Covered:

- **Core Python Functions:**

- def syntax, return statements, lambda expressions
- *args and **kwargs for flexible APIs
- Use-cases in ML pipelines for utility or metrics functions

- **OOP Concepts in ML:**

- Class creation with __init__
- Instance vs class variables (e.g., for shared configs)
- Inheritance and overriding methods (e.g., base Trainer with CNN/RandomForest variants)
- Reusable design for ML code using class-based components (ModelTrainer, Logger, ConfigManager)

- **Decorators in MLOps:**

- Logging training start-end time
- Custom decorators for timing and exception handling
- How to wrap functionality across pipelines

- **Modular Code Organization:**

- src/, utils/, configs/, models/, logs/
- Structuring reusable ML components for collaboration
- Using __init__.py, and best practices for imports

- **File Handling:**

- Create directories using os.makedirs() and pathlib.Path()
- Check file existence, write logs/artifacts/model weights programmatically

1 Python and MLOps Foundations

- **Logging for Observability:**

- Set up logger with formatters and handlers
- Log at multiple levels: INFO, DEBUG, ERROR
- Different logs for train.log, error.log, experiment.log

- **Command-Line Interfaces:**

- Use argparse to configure ML scripts dynamically (e.g., --lr, --epochs, --model_path)
- Simulate real-world ML pipelines with parameter-driven runs



1 Python and MLOps Foundations

Session 2: NumPy & Pandas Essentials (Beginner)

Goal: Build expertise in handling and transforming numerical/tabular data using NumPy and Pandas, forming the foundation of data pipelines in MLOps.

Topics Covered:

- **Core Python Functions:**

- def syntax, return statements, lambda expressions
- *args and **kwargs for flexible APIs
- Use-cases in ML pipelines for utility or metrics functions

- **NumPy for Matrix Ops & Preprocessing:**

- Create structured arrays for input pipelines
- Slice, reshape, and broadcast input tensors (features)
- Perform linear algebra: dot products, inverses, SVD
- Use-case: prepare data for ML model input

- **Pandas for Tabular ML Data:**

- Load, explore, and clean CSV/Excel data
- Filter/select with loc, iloc, conditionals
- Treat missing data using .fillna(), .dropna()
- Encode categorical variables (get_dummies(), LabelEncoder)
- Perform aggregations and grouped operations for feature engineering

- **Memory & Pipeline Best Practices:**

- Efficient DataFrame chaining
- Avoid unnecessary .copy() calls and in-place mutations
- Use type conversion to reduce memory usage

1 Python and MLOps Foundations

Session 3: MLOps Overview & ML System Design (Beginner)

Goal: Grasp the MLOps lifecycle from development to production, understanding architecture design, tool selection, and team roles in real-world AI systems.

Topics Covered:

- **MLOps Fundamentals:**

- What is MLOps? Where does it differ from DevOps?
- Common problems solved by MLOps: reproducibility, automation, monitoring

- **ML System Lifecycle:**

- Data collection → Cleaning → Experimentation → Training → Evaluation → Deployment → Monitoring → Retraining

- **Tooling Landscape:**

- Git, MLflow, Docker, Kubernetes, Airflow, Prometheus, Grafana
- Importance of tracking: data versioning, code changes, experiment results

- **Architectural Styles:**

- Monolithic training scripts vs modular pipelines
- Batch inference vs real-time model serving
- Design patterns: training/inference segregation, microservices

- **Roles in an MLOps Team:**

- ML Engineer, MLOps Engineer, Data Scientist, Platform Engineer

1 Python and MLOps Foundations

Session 4: Git & CI/CD for ML (Intermediate)

Goal: Learn the principles of version control and build automation pipelines tailored to ML workflows using GitHub Actions.

Topics Covered:

- **Git Fundamentals for ML:**
 - Track notebooks, models, and datasets using .gitignore, .gitattributes
 - Create branches for model experimentation
 - Use tags for tracking production model versions
- **GitHub Collaboration:**
 - Pull requests and code reviews for ML code
 - Manage issues, enhancements, and release cycles
- **Git LFS:**
 - Track large files (datasets, model weights) efficiently
 - Configure .gitattributes and remote storage for ML artifacts
- **CI/CD Concepts for ML:**
 - Continuous Integration: test code, verify environments, linting
 - Continuous Delivery: deploy models, update inference endpoints
- **GitHub Actions Deep Dive:**
 - Define workflows to install Python env, run tests, lint code
 - Add stages: build, train, test, validate

1 Python and MLOps Foundations

Session 5: Model Tracking with MLflow (Intermediate)

Goal: Enable reproducibility and experiment comparison by using MLflow to log, organize, and visualize multiple ML model training runs.

Topics Covered:

- **Why Model Tracking Matters:**
 - Track different experiment runs and hyperparameters
 - Ensure reproducibility and enable collaboration
 - Compare performance metrics across runs
- **Introduction to MLflow:**
 - Components overview:
 - **Tracking:** Logging metrics, params, artifacts
 - **Projects:** Packaging ML code for reproducibility
 - **Models:** Model abstraction and deployment formats
 - **Registry:** Lifecycle and stage control of models
- **MLflow Tracking API:**
 - Use `mlflow.start_run()` to begin experiment tracking
 - Log:
 - Parameters (`mlflow.log_param`)
 - Metrics (`mlflow.log_metric`)
 - Artifacts (`mlflow.log_artifact`)
 - Models (`mlflow.sklearn.log_model`)
 - Use `autolog()` with popular libraries like `sklearn`, `keras`
- **MLflow UI:**
 - Launch and navigate the MLflow dashboard
 - Filter, compare, and visualize multiple experiment runs
 - Add tags and notes for better organization
- **Use-Cases:**
 - Comparing hyperparameters (e.g., learning rate, `max_depth`)
 - Logging dataset version used in training
 - Storing visualizations (e.g., ROC curve, confusion matrix) as artifacts

1 Python and MLOps Foundations

Session 6: MLflow Projects & Model Registry (Intermediate)

Goal: Standardize and operationalize ML training workflows with MLflow Projects and manage production-ready models with MLflow Model Registry.

Topics Covered:

- **MLflow Projects:**
 - Structure ML code with a MLproject file
 - Define entry_points, conda environments, and default parameters
 - Benefits for reproducibility, portability, and automation
 - Run project locally or remotely (e.g., MLflow server)
- **MLflow Model Registry:**
 - **Register Models** via code or UI
 - **Version Control**: Automatically increment versions
 - **Stage Transitions**:
 - None → Staging → Production → Archived
 - **Approval Workflows**:
 - Manually or programmatically promote/demote models
 - Attach metadata and comments (e.g., accuracy, reviewer)
- **Rollback and Reuse**:
 - Retrieve specific model versions
 - Revert to previous production-ready model

1 Python and MLOps Foundations

Session 7: Model Packaging and Versioning (Intermediate)

Goal: Learn how to serialize, version, and manage models for distribution, deployment, and rollback across environments.

Topics Covered:

- **Serialization Techniques:**

- Use pickle, joblib for Python-based models
- Use ONNX, TorchScript for framework-agnostic, portable deployment
- Considerations: file size, compatibility, security

- **Versioning Best Practices:**

- Follow Semantic Versioning (SemVer): MAJOR.MINOR.PATCH
- Git tags for syncing code versions with model versions
- Save metadata (e.g., training timestamp, dataset hash)

- **Reproducibility:**

- Save environment dependencies: requirements.txt, conda.yaml
- Log model signature (input/output schema)
- Capture random seeds and random states for full reproducibility

1 Python and MLOps Foundations

Session 8: End-to-End CI/CD for ML with DVC (Advanced)

Goal: Build a reproducible and automated ML pipeline using DVC for dataset and model versioning, integrated with GitHub Actions for CI/CD.

Topics Covered:

- **Serialization Techniques:**

- Use pickle, joblib for Python-based models
- Use ONNX, TorchScript for framework-agnostic, portable deployment
- Considerations: file size, compatibility, security

- **Versioning Best Practices:**

- Follow Semantic Versioning (SemVer): MAJOR.MINOR.PATCH
- Git tags for syncing code versions with model versions
- Save metadata (e.g., training timestamp, dataset hash)

- **Reproducibility:**

- Save environment dependencies: requirements.txt, conda.yaml
- Log model signature (input/output schema)
- Capture random seeds and random states for full reproducibility

2 Modern Cloud-Native MLOps

Session 9: Orchestration with Apache Airflow (Advanced)

Goal: Build scalable, repeatable ML pipelines using Airflow's DAG-based orchestration.

Topics Covered:

- **Why Workflow Orchestration in MLOps:**
 - Automate and schedule end-to-end ML workflows
 - Enable retry logic, failure notifications, and inter-task dependencies
 - Maintain pipeline reproducibility and auditability
- **Airflow Concepts:**
 - **DAG (Directed Acyclic Graph):** Defines a pipeline of tasks with dependencies
 - **Operators:** Define task logic (e.g., PythonOperator, BashOperator)
 - **Scheduler & Workers:** Schedule and execute tasks asynchronously
 - **Metadata DB:** Tracks DAG state and task history
- **Modular Pipeline Design:**
 - Break workflow into logical steps: data ingestion → preprocessing → training → evaluation → deployment
 - Use templated variables/macros for reusability
- **Airflow UI:**
 - Visualize DAG structure and execution flow
 - Monitor task retries, failures, and logs
 - Trigger DAGs manually or via schedules (cron)

2 Modern Cloud-Native MLOps

Session 10: Hyperparameter Tuning with Optuna (Intermediate)

Goal: Improve model accuracy and generalization by automatically finding optimal hyperparameters.

Topics Covered:

- **Importance of Hyperparameter Tuning:**
 - Manual tuning is inefficient and non-reproducible
 - Automated tuning ensures systematic exploration of parameter space
- **Optuna Architecture:**
 - **Study:** A full optimization session
 - **Trial:** A single evaluation run with a hyperparameter set
 - **Sampler:** Chooses next trial's parameters (e.g., TPE, Random)
 - **Pruner:** Stops unpromising trials early to save compute
- **Integration:**
 - Plug into Scikit-learn, PyTorch, XGBoost, LightGBM
 - Works with `train_test_split`, cross-validation, early stopping
- **Visualization:**
 - Use Optuna's dashboard or Matplotlib plots for:
 - Optimization history
 - Parameter importance
 - Slice plots
- **MLflow Logging:**
 - Combine Optuna with MLflow to persist trial metrics and artifacts

2 Modern Cloud-Native MLOps

Session 11: Kubeflow Pipelines (KFP) (Advanced)

Goal: Create scalable, portable ML pipelines using the Kubeflow SDK running on Kubernetes.

Topics Covered:

- **What is Kubeflow?**

- End-to-end ML toolkit built on Kubernetes
- Tailored for repeatable training, tuning, deployment

- **KFP vs Airflow:**

- Airflow is general-purpose; KFP is ML-specific
- KFP integrates natively with containerized ML workflows
- Use KFP when running on cloud-native infra with Kubernetes

- **Pipeline Components:**

- Each step is a Docker container
- Input/output artifacts are tracked automatically
- Use KFP DSL to define Pythonic workflows

- **Pipeline Compilation & Deployment:**

- Use `kfp.Client()` to compile and upload pipelines
- Trigger experiments, monitor via UI, log metrics

2 Modern Cloud-Native MLOps

Session 12: Kubernetes for MLOps (Advanced)

Goal: Understand how Kubernetes manages ML infrastructure and deployment lifecycle.

Topics Covered:

- **Why Kubernetes?**

- Automates scaling, resilience, and resource utilization
- Consistent environment for dev → staging → prod

- **Kubernetes Core Concepts:**

- **Pods:** Smallest deployable unit, encapsulates containers
- **Deployments:** Manages Pod replication and rolling updates
- **Services:** Exposes Pods to internal/external network
- **ConfigMaps & Secrets:** Manage configuration securely
- **Namespaces:** Isolate environments (e.g., staging vs prod)

- **Kubernetes for ML:**

- Deploying models as RESTful microservices
- Define resource limits for compute-bound workloads
- Use Horizontal Pod Autoscaler (HPA)

2 Modern Cloud-Native MLOps

Session 13: Model Deployment with Kubernetes & TF Serving (Advanced)

Goal: Serve ML models efficiently using Kubernetes-native serving tools.

Topics Covered:

- **Serving Strategies:**
 - Batch vs real-time inference
 - REST vs gRPC API protocols
- **TensorFlow Serving:**
 - Export SavedModel format
 - Serve multiple model versions with auto-routing
 - Supports batching, latency control
- **Kubernetes + Model Serving:**
 - Use InferenceService from KServe or custom Deployment
 - Deploy TF Serving or custom container
 - Use YAML to define autoscaling, resource quotas
- **Custom Python Predictors:**
 - Write predict() functions with preprocessing/postprocessing
 - Bundle model and logic in container

2 Modern Cloud-Native MLOps

Session 14: Monitoring with Prometheus & Grafana (Intermediate)

Goal: Monitor models and infrastructure for performance, reliability, and drift detection.

Topics Covered:

- **What to Monitor in MLOps:**
 - Model: latency, throughput, error rate, drift
 - Infra: CPU/GPU utilization, memory, disk, network
- **Prometheus Overview:**
 - Pulls metrics from /metrics endpoint
 - Stores time-series data for querying and alerting
- **Grafana:**
 - Visualize metrics with customizable dashboards
 - Integrate alerts (Slack, email, PagerDuty)
- **Integration with ML APIs:**
 - Use Prometheus Python client to expose metrics
 - Monitor business-level KPIs: avg response time, failure count

2 Modern Cloud-Native MLOps

Session 15: Advanced Deployment Patterns (Intermediate)

Goal: Release models safely in production using advanced deployment strategies and observability.

Topics Covered:

- **Deployment Strategies:**
 - **Blue-Green:** Deploy two environments, switch traffic
 - **Canary:** Gradual rollout to subset of users (e.g., 20% first)
 - **Shadow Testing:** Send traffic to new model but don't expose results
- **Traffic Control in Kubernetes:**
 - Use Ingress controllers or KServe traffic splits
 - Automate traffic shifting based on time or metrics
- **Rollbacks:**
 - Track model version and metrics
 - Revert automatically if thresholds are breached

2 Modern Cloud-Native MLOps

Session 16: Feedback Loops & Cost Optimization (Intermediate)

Goal: Build self-learning pipelines by incorporating feedback, and manage cloud cost efficiently.

Topics Covered:

- **Human-in-the-Loop Feedback:**
 - Log predictions + user feedback (e.g., “was this correct?”)
 - Store in a feature store or feedback table
- **Drift & Retraining Triggers:**
 - Data drift: change in input distributions
 - Concept drift: change in label relationship
 - Use Kolmogorov-Smirnov tests, population stability index
- **Cost Optimization Techniques:**
 - Use spot/preemptible instances for training
 - Schedule batch inference instead of real-time
 - Apply model quantization to reduce resource usage

3 Cloud with Productionization

Session 17: AWS ML Stack Overview (Intermediate)

Goal: Get hands-on with AWS-native ML tools to build automated, event-driven pipelines from data ingestion to model inference.

Topics Covered:

- **AWS ML Ecosystem Overview:**
 - **Amazon SageMaker:** End-to-end ML lifecycle platform (train, deploy, monitor, pipeline)
 - **AWS Lambda:** Serverless compute triggered by events (e.g., new S3 file)
 - **Amazon S3:** Object storage for datasets and model artifacts
 - **Amazon ECR:** Container registry for custom Docker images
 - **CloudWatch:** Logging and monitoring for serverless and ML workloads
- **Event-Driven ML Pipelines:**
 - Real-world scenario: A new file in S3 triggers a Lambda function
 - Lambda parses file → launches SageMaker job → saves result → logs to CloudWatch
 - No need for managing servers or cron jobs

3 Cloud with Productionization

Session 18: SageMaker Pipelines with Model Registry (Advanced)

Goal: Automate and manage the entire ML lifecycle with SageMaker Pipelines and register models using built-in model registry.

Topics Covered:

- **SageMaker Pipeline Components:**
 - **ProcessingStep:** For data cleaning, feature engineering
 - **TrainingStep:** Launch ML training jobs
 - **ModelStep:** Create model from training artifacts
 - **ConditionStep:** Conditional logic based on metric thresholds
 - **RegisterModel:** Store trained model with metadata
- **Pipeline Building with Python SDK:**
 - Define parameters (e.g., learning_rate)
 - Chain steps and define dependencies
 - Visualize pipeline execution in SageMaker Studio
- **Model Registry:**
 - Register trained model to track versions
 - Transition models from "None" → "Staging" → "Production"
 - Store metadata like metrics, training job ID, deployment URL

3 Cloud with Productionization

Session 19: Drift Monitoring in Production (Advanced)

Goal: Detect when your production model starts to degrade due to changes in data or concept.

Topics Covered:

- **Types of Drift:**
 - **Data Drift:** Input distribution has changed (e.g., new age group in user data)
 - **Concept Drift:** The relationship between features and labels has changed (e.g., sentiment meaning changes over time)
 - **Label Drift:** Class distribution changes in output (e.g., more fraudulent transactions)
- **Drift Detection Techniques:**
 - **KS Test:** Statistical test for comparing distributions
 - **Population Stability Index (PSI):** Measures how stable a variable's distribution is over time
 - **Model Score Drift:** Distribution of predicted scores changing
- **SageMaker Model Monitor:**
 - Automatically collects endpoint invocations
 - Compares them against a **baseline** dataset
 - Can log alerts, metrics, and trigger retraining pipelines

3 Cloud with Productionization

Session 20: Multi-Cloud MLOps Integration (Advanced)

Goal: Understand how to design cloud-agnostic ML pipelines that combine tools across AWS, GCP, and Azure while maintaining governance and scalability.

Topics Covered:

- **Challenges in Multi-Cloud MLOps:**
 - Data silos and inconsistent APIs
 - Governance and security across platforms
 - Managing costs and compliance in regulated industries
- **Service Comparison:**
 - **SageMaker (AWS) vs Vertex AI (GCP) vs Azure ML**
 - Differences in training orchestration, model registry, pipelines, SDKs
 - When to prefer one over another
- **Cloud-Agnostic Pipelines:**
 - Use open tools like **MLflow, KServe, Kubernetes, DVC**
 - Abstract storage with S3, GCS, or MinIO
 - Use Terraform to define infra across clouds
- **Hybrid Strategy Example:**
 - Train model on GCP Vertex AI
 - Register model in MLflow
 - Deploy on AWS using SageMaker or EKS

3 Cloud with Productionization

Sessions 21 & 22: Capstone Project 1 - End-to-End MLOps Pipeline (Advanced)

Goal: Build a full production-grade MLOps pipeline integrating tools from previous sessions, simulating a real business problem.

Topics Covered:

- **Use-Case Definition:**

- Choose a real-world use case (e.g., loan default prediction, email classification, review sentiment, product recommendations)
- Decide problem type: regression, classification, NLP

- **Architecture Blueprint:**

- **Data ingestion:** From CSV, API, or database
- **Preprocessing:** Transform, clean, validate data
- **Model training:** Train with versioned code and data
- **Experiment tracking:** MLflow logging of metrics/artifacts
- **CI/CD pipeline:** Automate testing, packaging, and deployment
- **Model Registry:** Store versions and manage promotion
- **Deployment:** Kubernetes + InferenceService (or SageMaker endpoint)

Monitoring: Track latency, errors, and drift with Prometheus/Grafana or Model Monitor

4 Agentic AI & LLM Ops

Session 23: LLM Lifecycle + Prompt Engineering (Intermediate)

Goal: Learn how to design effective prompts and understand the lifecycle of LLM applications.

Topics Covered:

- **LLM Lifecycle Overview:**
 - **Pretraining:** Models trained on large corpora (e.g., GPT, LLaMA)
 - **Fine-tuning:** Specialized domain training
 - **Inference:** Deployment in real-world apps
- **Prompt Engineering:**
 - **Zero-shot:** Ask directly without examples
 - **Few-shot:** Provide 1–3 examples to guide generation
 - **Chain-of-thought (CoT):** Guide reasoning through intermediate steps
- **Prompt Templates:**
 - Use Jinja-style format in LangChain for reusable prompts
 - Introduce **guardrails** to filter or sanitize output
- **LangChain Prompt Utilities:**
 - PromptTemplate, ChatPromptTemplate
 - Manage dynamic inputs, temperature, roles

4 Agentic AI & LLMOps

Session 24: Fine-tuning vs Prompt-tuning (Advanced)

Goal: Compare cost-effective LLM adaptation strategies, and fine-tune a small model using LoRA or QLoRA.

Topics Covered:

- **Adaptation Strategies:**

- **Full Fine-Tuning:** Update all weights – high cost, large gains
- **PEFT (Parameter-Efficient Fine-Tuning):** Modify adapters (LoRA), saves memory

- **Trade-Offs:**

- Inference latency, storage size, data sensitivity
- Quality vs cost for large vs small models

- **Tools & Frameworks:**

- Hugging Face transformers, peft, accelerate
- Google Colab with T4/A100 GPUs

4 Agentic AI & LLM Ops

Session 25: RAG Architecture with Vector DBs (Advanced)

Goal: Build scalable RAG (Retrieval-Augmented Generation) pipelines with embedding-based search.

Topics Covered:

- **What is RAG?**
 - Combine a retrieval step with an LLM generation step
 - Reduce hallucinations, ground responses in documents
- **Chunking & Embedding:**
 - Chunk long docs into semantically meaningful blocks
 - Convert to vectors using OpenAI, Cohere, or HuggingFace embeddings
- **Vector Databases:**
 - **FAISS:** Local, fast, easy to use
 - **Pinecone, Weaviate:** Scalable, managed
 - Indexing strategies and search accuracy
- **LangChain Integration:**
 - RetrievalQA, VectorStoreRetriever, ConversationalRetrievalChain

4 Agentic AI & LLMOps

Session 26: LLM Caching & Routing Techniques (Advanced)

Goal: Optimize system performance and cost by using caching and routing between models.

Topics Covered:

- **Caching Techniques:**
 - Cache common prompts + responses to speed up
 - Store embeddings or final completions (e.g., Redis, file, memory)
- **Routing:**
 - Use fallback models: expensive → cheaper
 - Context-aware routing based on user, input length, or budget
- **LangChain Tools:**
 - MultiPromptChain, LLMRouter, Cache, RetrieverRouter

4 Agentic AI & LLMOps

Session 27: Agentic Frameworks – LangChain, CrewAI, LangGraph (Advanced)

Goal: Build systems with autonomous, role-based agents that collaborate to complete tasks.

Topics Covered:

- **Agentic AI Explained:**
 - Agents have tools, memory, and goals
 - Act autonomously with decision-making ability
- **LangChain Agents:**
 - Tool usage, memory integration (e.g., search, calculator)
- **CrewAI Framework:**
 - Define **Roles** (e.g., researcher, writer)
 - Build **Crew** of agents for collaborative tasks
- **LangGraph:**
 - Use graph-based workflows for deterministic execution

4 Agentic AI & LLM Ops

Session 28: Evaluation Tools – RAGAS, LangSmith, TruLens (Advanced)

Goal: Learn how to evaluate LLM app performance using open and proprietary evaluation frameworks.

Topics Covered:

- **Why LLM Evaluation is Hard:**
 - No one-size-fits-all metric
 - Output is variable and subjective
- **Metrics:**
 - **Faithfulness:** grounded in source
 - **Relevance:** is it useful/related?
 - **Conciseness, Toxicity, Hallucination rate**
- **Tools:**
 - **RAGAS:** Evaluate RAG pipelines
 - **LangSmith:** Track execution traces, latency, tool calls
 - **TruLens:** Quantitative feedback + human scoring

4 Agentic AI & LLMOps

Session 29: Evaluating LLMs – Metrics, Hallucination, Quality (Advanced)

Goal: Build LLM evaluation pipelines and integrate them with CI/CD for regression testing.

Topics Covered:

- **Types of Evaluation:**

- **Static:** Against ground truth labels
- **Dynamic:** Involving humans-in-the-loop
- **Hybrid:** Combine human feedback + auto-metrics

- **Tools:**

- **PromptLayer:** Prompt tracking and logging
- **Promptfoo:** Prompt testing toolkit
- **OpenAI Eval:** Model vs model comparison

- **Integration with CI/CD:**

- Regression tests for prompts
- Threshold-based alerting on quality drops

4 Agentic AI & LLMOps

Session 30: Capstone Project 2 – Build an Agentic AI System (Advanced)

Goal: Build and deploy a complete agentic AI system with RAG, vector DB, evaluation, and deployment.

Topics Covered:

- **Project Definition Examples:**

- Legal contract analyzer
- Sales email writer from CRM notes
- Support ticket summarizer with RAG + memory

- **Architecture:**

- Input → Vector Store → Retriever → Planner Agent → Tool → Summarizer Agent → Output
- Use LangChain or CrewAI for orchestration
- Deploy using LangServe or FastAPI

- **Evaluation Integration:**

- RAGAS for scoring
- LangSmith for pipeline trace logging
- TruLens for response quality rating

